

## Experiment 3

### Thread Creation and Execution<sup>1</sup>

#### Objectives

- ✓ To learn threads and multiprocessing (and multithreading). The primary objective of this lab is to implement the Thread Management Functions:
  - Creating Threads
  - Terminating Thread Execution
  - Passing Arguments To Threads
  - Thread Identifiers
  - Joining Threads
  - Detaching / Undetaching Threads

#### Prelab Activities

- ✓ Read the manual and try to do the experiment yourself before the lab.
- ✓ Write (copy and paste) and compile the codes given in document. Some programs may not be runnable. You should modify them according to the problem definitions given in the part of Exercises.

Note: You can simply ignore type-casting warnings.

#### General Information

##### A Thread:

***A thread is a semi-process, which has its own stack, and executes a given piece of code.***

Unlike a real process, the thread normally shares its memory with other threads (where as for processes we usually have a different memory area for each one of them). A Thread Group is a set of threads all executing inside the same process. They all share the same memory, and thus can access the same global variables, same heap memory, same set of file descriptors, etc. All these threads execute in parallel (i.e. using time slices, or if the system has several processors, then really in parallel).

##### pthread:

Historically, hardware vendors have implemented their own proprietary versions of threads. These implementations differed substantially from each other making it difficult for programmers to develop portable threaded applications.

In order to take full advantage of the capabilities provided by threads, a standardized programming interface was required. For UNIX systems, this interface has been specified by the IEEE POSIX 1003.1c standard (1995). Implementations which adhere to this standard are referred to as POSIX threads, or Pthreads. Most hardware vendors now offer Pthreads in addition to their proprietary API's.

---

<sup>1</sup>

Ref: M. Akbar Badhusha, King Fahd University of Petroleum and Minerals, Saudi Arabia.

**Pthreads** are defined as a set of C language programming types and procedure calls. Vendors usually provide a Pthreads implementation in the form of a header/include file and a library which you link with your program. The primary motivation for using Pthreads is to realize potential program performance gains.

When compared to the cost of creating and managing a process, a thread can be created with much less operating system overhead. Managing threads requires fewer system resources than managing processes.

All threads within a process share the same address space. Inter-thread communication is more efficient and in many cases, easier to use than inter-process communication.

Threaded applications offer potential performance gains and practical advantages over non-threaded applications in several other ways:

**Overlapping CPU work with I/O:** For example, a program may have sections where it is performing a long I/O operation. While one thread is waiting for an I/O system call to complete, CPU intensive work can be performed by other threads. **Priority/real-time scheduling:** tasks which are more important can be scheduled to supersede or interrupt lower priority tasks. **Asynchronous event handling:** tasks which service events of indeterminate frequency and duration can be interleaved. For example, a web server can both transfer data from previous requests and manage the arrival of new requests.

Multi-threaded applications will work on a uniprocessor system, yet naturally take advantage of a multiprocessor system, without recompiling.

In a multiprocessor environment, the most important reason for using Pthreads is to take advantage of potential parallelism. This will be the focus of the remainder of this session.

### The pthreads API:

The subroutines which comprise the Pthreads API can be informally grouped into three major classes:

**Thread management:** The first class of functions work directly on threads - creating, detaching, joining, etc. They include functions to set/query thread attributes (joinable, scheduling etc.)

**Mutexes:** The second class of functions deals with a coarse type of synchronization, called a "mutex", which is an abbreviation for "mutual exclusion". Mutex functions provide for creating, destroying, locking and unlocking mutexes. They are also supplemented by mutex attribute functions that set or modify attributes associated with mutexes.

**Condition variables:** The third class of functions deal with a finer type of synchronization - based upon programmer specified conditions. This class includes functions to create, destroy, wait and signal based upon specified variable values. Functions to set/query condition variable attributes are also included.

**Naming conventions:** All identifiers in the threads library begin with pthread\_

pthread_	Threads themselves and miscellaneous subroutines
pthread_attr	Thread attributes objects
pthread_mutex	Mutexes
pthread_mutexattr	Mutex attributes objects.
pthread_cond	Condition variables
pthread_condattr	Condition attributes objects
pthread_key	Thread-specific data keys

**Thread Management Functions:**

The function ***pthread\_create*** is used to create a new thread, and the function ***pthread\_exit*** is used by a thread to terminate itself. The function ***pthread\_join*** is used by a thread to wait for termination of a thread.

<b>Function:</b>	<pre> Int pthread_create(     pthread_t *threadhandle,     /*Thread handle returned by reference*/     pthread_attr_t *attribute,     /*Special Attribute for starting thread, may be NULL */     void *(*start_routine)(void *),     /* Main Function which thread executes */     void *arg     /* An extra argument passed as a pointer */ ); </pre>
<b>Info:</b>	Request the <b>PThread</b> library for <b>creation</b> of a new thread. The return value is <b>0</b> on <b>success</b> . The return value is <b>negative</b> on <b>failure</b> . The <b>pthread_t</b> is an abstract datatype that is used as a handle to <b>reference</b> the thread.

<b>Function:</b>	<pre> void pthread_exit(     void *retval /* return value passed as a pointer */ ); </pre>
<b>Info:</b>	This Function is used by a thread to <b>terminate</b> . The return value is passed as a <b>pointer</b> . This pointer value can be anything so long as it does not exceed the size of (void *). Be careful, this is system dependent. You may wish to return an address of a structure, if the returned data is very large.

<b>Function:</b>	<pre> int pthread_join(     pthread_t threadhandle, /* Pass threadhandle */     void **returnvalue /*Return value is returned by ref.*/ ); </pre>
<b>Info:</b>	Return <b>0</b> on <b>success</b> , and <b>negative</b> on <b>failure</b> . The returned value is a <b>pointer</b> returned by <b>reference</b> . If you do not care about the return value, you can pass <b>NULL</b> for the second argument.

**Thread Initialization:**

- Include the pthread.h library : `#include <pthread.h>`
- Declare a variable of type pthread\_t : `pthread_t the_thread`
- When you compile, add -lpthread to the linker flags :  
`gcc threads.c -o threads -lpthread`

Initially, threads are created from within a process. Once created, threads are peers, and may create other threads. Note that an "initial thread" exists by default and is the thread which runs main.

**Terminating Thread Execution:**

```
int pthread_cancel (pthread_t thread)
```

***pthread\_cancel*** sends a cancellation request to the thread denoted by the thread argument. If there is no such thread, ***pthread\_cancel*** fails. Otherwise it returns 0.

A cancel is a mechanism by which a calling thread informs either itself or the called thread to terminate as quickly as possible. Issuing a cancel does not guarantee that the canceled thread receives or handles the cancel. The canceled thread can delay processing the cancel after receiving it. For instance, if a cancel arrives during an important operation, the canceled thread can continue if what it is doing cannot be interrupted at the point where the cancel is requested.

The programmer may specify a termination status, which is stored as a void pointer for any thread that may join the calling thread.

***There are several ways in which a pthread may be terminated:***

- The thread returns from its starting routine (the main routine for the initial thread). By default, the pthreads library will reclaim any system resources used by the thread. This is similar to a process terminating when it reaches the end of main.
- The thread makes a call to the ***pthread\_exit*** subroutine (covered below).
- The thread is canceled by another thread via the ***pthread\_cancel*** routine (not covered here).
- The thread receives a signal that terminates it
- The entire process is terminated due to a call to either the exec or exit subroutines.

**Thread Attributes:**

Threads have a number of attributes that may be set at creation time. This is done by filling a thread attribute object ***attr*** of type ***pthread\_attr\_t***, then passing it as second argument to ***pthread\_create***. Passing NULL is equivalent to passing a thread attribute object with all attributes set to their default values.

Attribute objects are consulted only when creating a new thread. The same attribute object can be used for creating several threads. Modifying an attribute object after a call to ***pthread\_create*** does not change the attributes of the thread previously created.

```
int pthread_attr_init (pthread_attr_t *attr)
```

***pthread\_attr\_init*** initializes the thread attribute object ***attr*** and fills it with default values for the attributes. Each attribute ***attrname*** can be individually set using the function ***pthread\_attr\_setattrname*** and retrieved using the function ***pthread\_attr\_getattrname***.

```
int pthread_attr_destroy (pthread_attr_t *attr)
```

***pthread\_attr\_destroy*** destroys the attribute object pointed to by ***attr*** releasing any resources associated with it. ***attr*** is left in an undefined state, and you must not use it again in a call to any pthreads function until it has been reinitialized.

```
int pthread_attr_setattr (pthread_attr_t *obj, int value)
```

Set attribute **attr** to value in the attribute object pointed to by obj. See below for a list of possible attributes and the values they can take. On success, these functions return 0.

```
int pthread_attr_getattr(const pthread_attr_t *obj, int *value)
```

Store the current setting of **attr** in **obj** into the variable pointed to by value. These functions always return 0.

The following thread attributes are supported:

**'detachstate':** Choose whether the thread is created in the joinable state (value PTHREAD\_CREATE\_JOINABLE) or in the detached state (PTHREAD\_CREATE\_DETACHED). The default is PTHREAD\_CREATE\_JOINABLE. In the joinable state, another thread can synchronize on the thread termination and recover its termination code using pthread\_join, but some of the thread resources are kept allocated after the thread terminates, and reclaimed only when another thread performs pthread\_join on that thread. In the detached state, the thread resources are immediately freed when it terminates, but pthread\_join cannot be used to synchronize on the thread termination. A thread created in the joinable state can later be put in the detached thread using pthread\_detach.

**'schedpolicy':** Select the scheduling policy for the thread: one of SCHED\_OTHER (regular, non-realtime scheduling), SCHED\_RR (realtime, round-robin) or SCHED\_FIFO (realtime, first-in first-out). The default is SCHED\_OTHER. The realtime scheduling policies SCHED\_RR and SCHED\_FIFO are available only to processes with superuser privileges. pthread\_attr\_setschedparam will fail and return ENOTSUP if you try to set a realtime policy when you are unprivileged. The scheduling policy of a thread can be changed after creation with pthread\_setschedparam.

**'schedparam':** Change the scheduling parameter (the scheduling priority) for the thread. The default is 0. This attribute is not significant if the scheduling policy is SCHED\_OTHER; it only matters for the realtime policies SCHED\_RR and SCHED\_FIFO. The scheduling priority of a thread can be changed after creation with pthread\_setschedparam.

Thread Identifiers:

```
pthread_self ( )
```

Returns the unique thread ID of the calling thread. The returned data object is opaque can not be easily inspected.

```
pthread_equal ( thread1, thread2 )
```

Compares two thread IDs: If the two IDs are different 0 is returned, otherwise a non-zero value is returned. Because thread IDs are opaque objects, the C language equivalence operator == should not be used to compare two thread IDs.

**Example: Pthread Creation and Termination:**

```
#include <stdio.h>
#include <pthread.h>

//Are the process id numbers of parent and child thread the same or different?
void *kidfnc(void *p){
    printf ("kidfnc ID is ---> %d\n", getpid());
}

void main (){
    pthread_t kid;

    pthread_create (&kid, NULL, kidfnc, NULL);
    printf ("Parent ID is ---> %d\n", getpid());
    pthread_join (kid, NULL);
    printf ("No more kid!\n");
}
```

**Program 1.**

Sample output:

```
Parent ID is ---> 10059
kidfnc ID is ---> 10059
No more kid!
```

```
#include <stdio.h>
#include <pthread.h>

//Do the threads have separate copies of glob_data?
//Execute this program multiple times!
int glob_data = 5;

void *kidfunc(void *p){
    printf("(Kid) Global data = %d\n", glob_data);
    printf("(Kid) Global data = %d\n", glob_data = 15);
}

void main(){
    pthread_t kid;
    pthread_create (&kid, NULL, kidfunc, NULL);
    printf("(Par) Global data = %d\n", glob_data);
    printf("(Par) Global data = %d\n", glob_data = 10);
    pthread_join (kid, NULL) ;
    printf("!END! Global data = %d\n", glob_data);
}
```

**Program 2.**

Sample output:

```
(Par) Global data = 5
(Par) Global data = 10
(Kid) Global data = 10
(Kid) Global data = 15
!END! Global data = 15
```

**Multiple Threads:**

The example code below creates 5 threads with the `pthread_create()` routine. Each thread prints a "Hello World!" message, and then terminates with a call to `pthread_exit()`.

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5

void *printHello(void *threadid){
    printf("\n%d: Hello World!\n", threadid);
    pthread_exit(NULL);
}

void main(){
    pthread_t threads[NUM_THREADS];
    int rc, t;

    for(t=0; t < NUM_THREADS; t++){
        printf ("Creating thread %d\n", t);
        rc=pthread_create(&threads[t],NULL,printHello, (void *)t);
        if(rc){
            printf("ERROR;return code from pthread_create() is %d\n",rc);
            exit(-1);
        }
    }
}
```

**Program 3.**

Sample output:

```
Creating thread 0
Creating thread 1
Creating thread 2

0: Hello World!
Creating thread 3

1: Hello World!
Creating thread 4

2: Hello World!

3: Hello World!

4: Hello World!
```

**Difference between process and threads :**

```
#include <stdio.h>
#include <pthread.h>
#include <unistd.h>

int dataGlobal;
void thread_func(void *ptr);

int main(){
    int dataLocal, pid, status;
    pthread_t thread1, thread2;

    printf("First, we create two threads to ");
    printf("see better what context they share...\n");

    dataGlobal=1000;
    printf("Set dataGlobal=%d\n",dataGlobal);
    pthread_create(&thread1,NULL, (void*)&thread_func, (void*) NULL);
    pthread_create(&thread2,NULL, (void*)&thread_func, (void*) NULL);

    pthread_join(thread1, NULL);
    pthread_join(thread2, NULL);
    printf("After threads, dataGlobal=%d\n",dataGlobal);

    printf("\nNow that the threads are done, let's call fork..\n");
    dataGlobal=dataLocal=17;
    printf("Before fork(), dataLocal=%d, ",dataLocal);
    printf("dataGlobal=%d\n",dataGlobal);

    pid=fork();
    if (0 == pid) { /* this is the child */
        printf("(child ) pid %d: &global: ", getpid());
        printf("%X, &local: %X\n", &dataGlobal, &dataLocal);

        dataLocal=13; dataGlobal=23; //new values
        printf("(child ) dataLocal=%d, ",dataLocal);
        printf("dataGlobal=%d\n",dataGlobal);
        exit(0);
    }
    else { /* this is the parent */
        printf("(parent) pid %d: &global: ", getpid());
        printf("%X, &local: %X\n", &dataGlobal, &dataLocal);

        wait(&status);
        printf("(parent) dataLocal=%d, ",dataLocal);
        printf("dataGlobal=%d\n",dataGlobal);
    }
    exit(0);
}

void thread_func(void *dummy)
{
    int local_thread;
    printf("Thread %u, pid %d, &global:%X, &local: %X\n", (unsigned
int)pthread_self(), getpid(),&dataGlobal, &local_thread);

    dataGlobal++;
    printf("Thread %u, incremented dataGlobal=%d\n", (unsigned
int)pthread_self(), dataGlobal);
    pthread_exit(0);
}
```

**Program 4.**



Sample output:

First, we create two threads to see better what context they share...

Set dataGlobal=1000

Thread 4187383552, pid 11749, &global:70777014, &local: F9965ED4

Thread 4187383552, incremented dataGlobal=1001

Thread 4195776256, pid 11749, &global:70777014, &local: FA166ED4

Thread 4195776256, incremented dataGlobal=1002

After threads, dataGlobal=1002

Now that the threads are done, let's call fork..

Before fork(), dataLocal=17, dataGlobal=17

(parent) pid 11749: &global: 70777014, &local: 54BF09DC

(child ) pid 11752: &global: 70777014, &local: 54BF09DC

(child ) dataLocal=13, dataGlobal=23

(parent) dataLocal=17, dataGlobal=17

```
#include <stdio.h>
#include <pthread.h>
#include <stdlib.h>

#define NKIDS 50
int tot_items = 0;
struct kidrec{
    int data;
    pthread_t id;
};

void *kidfunc(void *p){
    int *ip = (int *)p;
    int tmp, n;
    tmp = tot_items;
    for (n = 50000; n--;)
        tot_items = tmp + *ip;
}

void main(){
    struct kidrec kids[NKIDS];
    int m;

    for (m=0; m<NKIDS; ++m){
        kids[m].data = m+1 ;
        pthread_create(&kids[m].id, NULL, kidfunc, &kids[m].data);
    }
    for (m=0; m<NKIDS; ++m)
        pthread_join (kids[m].id, NULL);

    printf ("End! Grand Total = %d\n", tot_items);
}
```

**Program 5.**Sample output:

Run it several times until you see different output. How many times is the line **tot\_items = tmp + \*ip;** executed? What values does **\*ip** have during these executions?

**Passing Arguments to Threads :**

The ***pthread\_create()*** routine permits the programmer to pass one argument to the thread start routine. For cases where multiple arguments must be passed, this limitation is easily overcome by creating a structure which contains all of the arguments, and then passing a pointer to that structure in the ***pthread\_create()*** routine.

All arguments must be passed by reference and cast to (void \*).

Important: threads initially access their data structures in the parent thread's memory space. The data structure must not be corrupted/modified until thread has finished accessing it. The following example passes a simple integer to each thread.

**Example: pthread\_create() argument passing :**

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 7

char *messages[NUM_THREADS];

void *printHello(void *threadid){
    int *id_ptr, taskid;
    sleep(1);
    id_ptr = (int *) threadid;
    taskid = *id_ptr;
    printf("(from thread %d) %s\n", taskid, messages[taskid]);
    pthread_exit(NULL);
}

void main(){
    pthread_t threads[NUM_THREADS];
    int *taskids[NUM_THREADS], rc, t;

    messages[0] = "English: Hello World!";
    messages[1] = "French: Bonjour, le monde!";
    messages[2] = "Spanish: Hola al mundo";
    messages[3] = "Klingon: Nuq neH!";
    messages[4] = "German: Guten Tag, Welt!";
    messages[5] = "Russian: Zdravstvuyte, mir!";
    messages[6] = "Japan: Sekai e konnichiwa!";
    messages[7] = "Latin: Orbis, te saluto!";

    for(t=0;t<NUM_THREADS;t++){
        taskids[t] = (int *) malloc(sizeof(int));
        *taskids[t] = t;
        printf("Creating thread %d\n", t);
        rc=pthread_create(&threads[t],NULL,printHello,(void
*)taskids[t]);
        if (rc){
            printf("ERROR;return code of pthread_create() %d\n", rc);
            exit(-1);
        }
    }
    pthread_exit(NULL);
}
```

**Program 6.**

Sample output:

```
Creating thread 0
Creating thread 1
Creating thread 2
Creating thread 3
Creating thread 4
Creating thread 5
Creating thread 6
(from thread 0) English: Hello World!
(from thread 3) Klingon: Nuq neH!
(from thread 1) French:  Bonjour, le monde!
(from thread 6) Japan:  Sekai e konnichiwa!
(from thread 2) Spanish: Hola al mundo
(from thread 4) German:  Guten Tag, Welt!
(from thread 5) Russian: Zdravstvyyte, mir!
```

**Exercises:**

**Problem 1.** The following Program 7 demonstrates a simple program where the main thread creates another thread to print out the numbers from 1 to 20. The main thread waits till the child thread finishes.

```
#include <pthread.h>
#include <stdio.h>

void childFunc(void *argument){
    int i;
    for (i = 1; i <= 20; ++i ){
        printf(" Child Count - %d\n", i);
    }
    pthread_exit(0);
}

int main(void){
    pthread_t hThread;
    int ret;

    ret=pthread_create(&hThread,NULL,(void*)childFunc, NULL);

    if (ret < 0) {
        printf("Thread Creation Failed\n");
        return 1;
    }

    pthread_join (hThread, NULL); /* Parent waits for */
    printf("Parent is continuing....\n");
    return 0;
}
```

**Program 7.**

Compile and execute the Program 7. Show the output and explain why is the output so?

**Problem 2.** In the Program 8. modify the above Program 7. such that the main program passes the count as argument to the child thread function and the child thread function prints that many count print statements.

```
#include <pthread.h>
#include <stdio.h>

void childFunction (int argument)
{
    int i;

    //.....

    pthread_exit(0);
}

int main(void)
{
    pthread_t hThread;
    //pthread_create (.....);

    pthread_join (hThread, NULL);

    printf ("Parent is continuing....\n");
    return 0;
}
```

**Program 8.**

Modify, compile and execute the Program 8., and show the output and explain why is the output so?

**Problem 3.** Write a program Program 9. by removing pthread\_exit function from child thread function and check the output? Is it the same as output of Program 8.? If so Why? Explain?